

CC lab 2

Name: Arnav Sinha
SRN: PES2UG23CS092
Sec: B

SS1

The screenshot shows a web application interface for event registration. At the top, there is a header bar with the text "Fest Monolith" and "FastAPI • SQLite • Locust". It also displays the user information "Logged in as PES2UG23CS092" and navigation links for "Events", "My Events", "Checkout", and "Logout".

The main content area is titled "Events" and displays a list of nine events:

- Event ID: 1** - **Hackathon**: ₹ 500. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 2** - **Dance**: ₹ 300. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 3** - **Hackathon**: ₹ 500. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 4** - **Dance Battle**: ₹ 300. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 5** - **AI Workshop**: ₹ 400. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 6** - **Photography Walk**: ₹ 200. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 7** - **Gaming Tournament**: ₹ 350. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 8** - **Music Night**: ₹ 250. Includes certificate • instant registration • limited seats. **Register**
- Event ID: 9** - **Treasure Hunt**: ₹ 150. Includes certificate • instant registration • limited seats. **Register**

A "View My Events" button is located in the top right corner of the event list.

SS2

The screenshot shows a web browser window with the URL `localhost:8000/checkout`. The page title is "Fest Monolith" and it displays a "Monolith Failure" message. The message states: "One bug in one module impacted the **entire application**". Below this, an "Error Message" is shown: "division by zero". To the right, a section titled "What should you do in the lab?" lists three steps: "Take a screenshot (crash demonstration)", "Fix the bug in the indicated module", and "Restart the server and verify recovery". At the bottom of the page, there is a log output:

```
INFO: Waiting for application startup.  
INFO: Application startup complete.  
INFO: 127.0.0.1:54038 - "GET /checkout HTTP/1.1" 500 Internal Server Error  
ERROR: Exception in ASGI application
```

SS3

The screenshot shows a web browser window with the URL `localhost:8000/checkout`. The page title is "Fest Monolith" and it displays a "Checkout" message. It states: "This route is used to demonstrate a monolith crash + optimization." A "Total Payable" amount of "₹ 6600" is shown. Below this, a note says: "After fixing + optimizing checkout logic, re-run Locust and compare results." To the right, a section titled "What you should observe" lists three points: "One buggy feature can crash the entire monolith.", "Inefficient loops cause high response times under load.", and "Optimization improves performance but architecture still scales as one unit.". A yellow box at the bottom right contains the text: "Next Lab: Split this monolith into Microservices (Events / Registration / Checkout)." At the bottom of the page, there is a log output:

```
INFO: Started server process [15172]  
INFO: Waiting for application startup.  
INFO: Application startup complete.  
INFO: 127.0.0.1:56450 - "GET /checkout HTTP/1.1" 200 OK
```

0 ▲ 0 Ln 10, Col 6 Spaces: 4 UTF-8

SS4

The screenshot shows a development environment with multiple windows. On the left is a Locust test results interface. The main table shows the following data:

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s	
GET	/checkout	20	0	10	2200	2200	124.12	4	2216	2797	0.7	0
	Aggregated	20	0	10	2200	2200	124.12	4	2216	2797	0.7	0

On the right, there is a code editor showing Python code for a file named `main.py`:

```

main.py  _init_.py  checkout_locustfile.py
CC Lab-2 > checkout > _init_.py > checkout_logic
3 def checkout_logic():
6
7     events = db.execute("SELECT fee FROM events").fetchall()
8
9     # Uncomment this line initially for the crash screenshot task
10    #! / #
11
12    total = 0
13    for e in events:
14        fee = e[0]
15        while fee > 0:
16            total += 1
17            fee -= 1
18
19    return total
20

```

Below the code editor is a terminal window showing command-line output. At the bottom, a PowerShell window shows the current directory as `C:\Users\arnav\Desktop\PE52UG23CS092\CC Lab-2>`.

SS5

The screenshot shows a development environment with multiple windows. On the left is a Locust test results interface. The main table shows the following data:

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s	
GET	/checkout	19	0	8	2100	2100	119.03	4	2099	2797	0.6	0
	Aggregated	19	0	8	2100	2100	119.03	4	2099	2797	0.6	0

On the right, there is a code editor showing Python code for a file named `main.py`, which is identical to SS4.

Below the code editor is a terminal window showing command-line output. It includes log entries from the Locust application and a detailed response time percentile chart. At the bottom, a PowerShell window shows the current directory as `C:\Users\arnav\Desktop\PE52UG23CS092\CC Lab-2>`.

Before Optimization

- Requests/sec (RPS): ~0.70**
- Average Response Time: ~124 ms**

After Optimization

- Requests/sec (RPS): ~0.64**
- Average Response Time: ~119 ms**

SS6

SS7

SS8

The screenshot shows a browser window displaying Locust test results and an adjacent code editor window.

Locust Test Results:

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s	
GET	/my-events?user=locust_user	19	0	46	2200	2200	160.99	32	2159	3144	0.7	0
	Aggregated	19	0	46	2200	2200	160.99	32	2159	3144	0.7	0

Code Editor (main.py):

```
from locust import HttpUser, task, between

class TestJourneyUser(HttpUser):
    wait_time = between(1, 2)

    @task
    def journey_locustfile(self):
```

Code Editor (journey_locustfile.py):

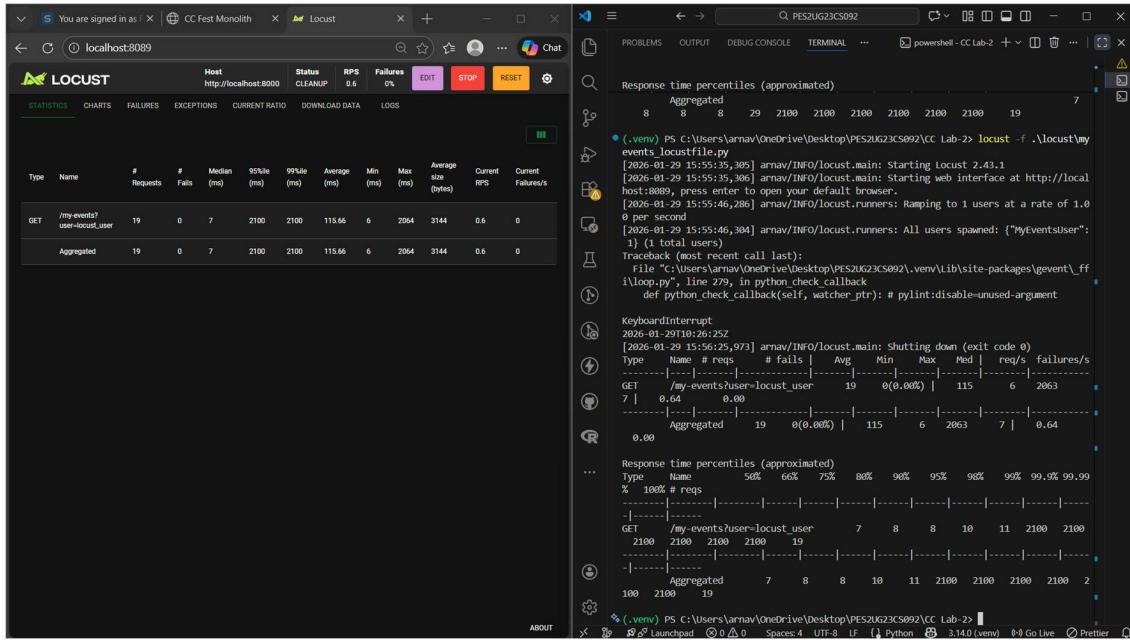
```
from locust import HttpUser, TaskSet, between

class MyTaskSet(TaskSet):
    @task
    def my_task(self):
```

IDE Terminal:

```
(venv) PS C:\Users\arnav\OneDrive\Desktop\PES2UG23CS092\CC Lab-2> locust -f .\locust\myevents\locustfile.py
[2026-01-29 15:23:20,668] arnav/INFO/locust.runners: All users spawned: {"MyEventsUser": 1} (1 total users)
Traceback (most recent call last):
  File "C:\Users\arnav\OneDrive\Desktop\PES2UG23CS092\.venv\lib\site-packages\gevent\ff
i_loop.py", line 279, in python_check_callback
    def python_check_callback(self, watcher_ptr): # pylint:disable=unused-argument
```

SS9



/events

- Bottleneck: Server-side latency on /events (DB/cache/query work), not the Locust client. With a 1–2s wait you're issuing low RPS, so any high p95/p99 tails you see are the backend's.
- Change you have now: Simple GET with query string in locust/events_locustfile.py, using the default 1–2s wait. Commented the line in main.py

```

64
65      # waste = 0
66      # for i in range(3000000):
67      #     waste += i % 3
68

```

- Why performance improves (when it does): Stability comes from low load (long waits), which reduces contention on the backend and hides bottlenecks; it's not a client-side optimization. If tails remain high, focus on backend profiling or caching.

/my-events

- Bottleneck: Same story—backend work for /my-events (auth/session lookups, DB reads).

- Change you have now: Simple GET with query string in locust/myevents_locustfile.py, 1–2s wait. Commented the line in main.py

```
# dummy = 0
# for _ in range(1500000):
#     dummy += 1
```

- Why performance improves (when it does): Lower request rate eases server pressure, so averages/tails can look better. Any remaining slowness points to the backend; you'd need shorter waits or more users to surface true capacity limits, then optimize the service (indexes, caching, N+1 fixes).

Repo: <https://github.com/arnavsinha20/CC-lab2>

Optimization:

- Short waits mean higher RPS so you can see how the backend behaves under load instead of idling between calls.
- Named requests and passing query params separately keep Locust stats clean and comparable.
- Timeouts and catch_response surface bad statuses quickly instead of masking errors.